Advanced Unix System Administration

Lecture 22 December 8, 2008

Steven Luo <sluo+decal@OCF.Berkeley.EDU>

- What is crypto?
 - Most prominently, set of algorithms, but these algorithms are bundled in a system!
- What can crypto do for you?
 - Confidentiality ensures that other people can't read your data
 - Authentication ensures that a message comes from its claimed source
 - Integrity ensures that data hasn't been modified or corrupted

- What does crypto NOT do for you?
 - Address unrelated weaknesses in your security model
 - Fix weaknesses in your implementation of a cryptosystem
 - Prevent user errors
 - Make users actually use your security system!
- Cryptography is NOT a magic bullet!
 - Crypto is well-studied; it's usually easier to find another way in

- Cryptographic primitives
 - Ciphers
 - Makes data unreadable except by holder(s) of a certain secret (the secret part is important!)
 - Symmetric ciphers share a secret for encryption and decryption, public-key ciphers only use a secret at decryption time
 - Public-key ciphers: RSA, ElGamal
 - Symmetric ciphers: DES, AES, RC4, ...
 - Block ciphers work on chunks of data, stream ciphers work on individual bytes
 - Block ciphers: DES, CAST family, Blowfish/Twofish, AES
 - Stream ciphers: RC4

- Cryptographic primitives con't
 - One-way hash functions
 - Take data and produce some short stream of bytes which somehow "identifies" that data
 - Important properties:
 - Small changes in the input should produce large changes in the output (avalanche effect)
 - Collisions should be difficult to generate
 - Should be difficult to calculate the original datastream from the hash
 - Not as well-studied as ciphers
 - Examples: MD5, RIPEMD-160, SHA family
 - NIST competition for next-gen hash ongoing

- Applications of crypto
 - Encrypting data
 - Apply a cipher to the data
 - Doesn't provide for integrity!
 - Data verification
 - Compare the hash of the data with a known hash
 - Where does one get the hash from?
 - Digital signatures
 - Concept: encrypt something with your private key, so that people can identify it as coming from you

- Applications of crypto con't
 - Hybrid cryptosystem
 - Hash the data, sign the hash, encrypt the data
 - Provides confidentiality, authentication, and integrity verification
 - Communications security
 - In general, you want to verify who you're talking to before you talk to them, so verify their public key
 - Has to be some infrastructure to verify this key!
 - Public-key algorithms are slow, so select a symmetric key via the PK-secured channel and switch to the symmetric cipher

- Applications of crypto con't
 - Secure authentication
 - Lots of possible schemes:
 - Require the user to sign some data of your choosing
 - Use hash functions on a secret
 - Send a secret over an encrypted channel
 - Carelessly designed schemes can be quite insecure!
 - Just accepting the hash of a secret alone leaves you open to replay attacks
 - What happens if one of the endpoints is malicious?

- Attacks on crypto
 - Cipher attacks
 - Data recovery: ability to read a message faster than brute force
 - Key recovery: ability to find the key used from the encrypted messages alone
 - Hash function attacks
 - Collision: generate two reasonably-related things that hash to the same value
 - Reversing the hash: find possible inputs to the hash from the output alone

- Attacks on crypto con't
 - Brute force
 - When your secrets are too small, it's possible to reverse a computation in a "reasonable" amount of time
 - What values provide "enough" security?
 - Symmetric ciphers: 128 bits
 - Public-key ciphers: 2048 bits
 - Hash functions: 256 bits
 - Note that computers and attacks improve with time!
 - Side-channel attacks
 - Lots of creative ways: measure power draw, time to run algorithm, etc.

- What do directory services do?
 - Classically, a directory service stores information about users and computers on a network
 - Internal address/telephone/email directories were some of the earliest applications for LDAP
 - In a system administration world, you frequently hear about it in the context of providing name and address lookups
 - Directories are optimized for fast read access via the network - hence the suitability

- Network Information Service (NIS/YP)
 - Another Sun RPC based service
 - Simple request-reply protocol, types of data available limited (passwd, hosts, etc.)
 - A limited generic "map" facility is available
 - Data is stored in flat files, with separate flat files needed for each field which can be looked up
 - No over-the-wire security

- NIS+
 - Looks sort of like NIS, but isn't NIS
 - Supports arbitrary key lookups and (weak) over-the-wire security
 - Also allows arbitrary data to be stored
 - Not known for its reliability and support
- Netware Directory Services, Windows NT domains
 - Similar design and scope as NIS

- Lightweight Directory Access Protocol
 - Originally intended as a lightweight, non-OSIstack based method of accessing X.500 directories
 - Protocol quickly was implemented standalone
 - Can store any type of data, in any organization
 - The data types and the structure are defined in "schemas"
 - With flexibility comes complexity
 - Lots of implementations

- LDAP con't
 - Lifecycle of an LDAP connection
 - The client "binds" to the server, authenticating and negotiating protocol version
 - Stream of requests issued
 - Various search operations and compare operations –
 LDAP mandates powerful built-in filters
 - Add, modify, delete entries
 - New operations can be defined via the "extended operation" operation
 - "Unbind" is the connection close
 - The use of TCP imposes some overhead

- LDAP con't
 - LDAP security
 - LDAP itself provides no over-the-wire security; connection security is usually managed via SSL/TLS
 - Secure authentication methods are provided by SASL, if both client and server support it
 - Use for Name Service Switch lookups
 - RFC 2307 (and later drafts) defines a standard schema for storage of this data that fits in with the standard schemas for other uses of LDAP
 - Overhead is considerable use a caching service!

- Active Directory
 - Embrace-extended version of LDAP and Kerberos
 - Standard (if buggy) LDAP and standard Kerberos but Microsoft uses a proprietary Kerberos-over-LDAP protocol to provide security
 - Can be used to store standard LDAP schema data, as well as M\$-proprietary schema data
 - Allows storing client policy and configuration
 - Interoperability with Unix can be difficult, but is possible

- Two forms of network storage:
 - Disk-level access give each client access to the underlying physical disk
 - Clients do reads from the disk and have to understand the filesystem themselves
 - Server can be simple, but more network I/O
 - SAN implementations (Fiber Channel, iSCSI)
 - File-level access clients operate on files
 - Clients request files, server fulfills requests
 - Server is more complex, saves on network I/O
 - NFS, SMB/CIFS, AFS/Coda ...

- The Network File System (NFS)
 - First developed at Sun in the 1980s, now controlled by IETF
 - Versions 2 and 3 attempt to provide (mostly) stateless operation, to simplify crash recovery
 - Features requiring state (locking) are implemented in additional protocols
 - Attempts to provide POSIX-like semantics within the possibilities of stateless operation
 - Note that NFS is NOT a POSIX-compliant filesystem!

- The Network File System (NFS) con't
 - Lifecycle of an NFS request
 - Client requests a "file handle" from the NFS mountd
 - Client requests that the nfsd perform operations on this file handle
 - Operations should be simple, atomic, and idempotent
 - In reality, they're not always
 - nfsd returns only when it has finished the operation or the operation has failed

- The Network File System (NFS) con't
 - Problems (in traditional NFS)
 - It's difficult to guarantee that repeating operations will be safe
 - Some features of POSIX filesystems, especially those requiring atomicity, are impossible to implement
 - The default mode can leave clients hanging for a very, very long time on a server crash
 - Security depends on the client saying who it is

- Other network file systems
 - As a rule, all maintain state
 - Allows for more complex commands, more featureful semantics, but makes crash recovery difficult
 - Most implement richer access control
 - The exact model tends to differ AFS, CIFS, and NFSv4 offer separate, differing access control models
 - Over-the-wire security may also be present
 - Some provide for distributed operation

iSCSI

- A translation of the SCSI command set to a network
 - SCSI topology already looks a bit like a network, so this isn't too hard
 - Reliability questions need to be dealt with
- An "initiator" on the host connects to storage on a "target"
- Simultaneous access from multiple initiators requires filesystem support (Red Hat GFS, OCFS)

- A brief history of Internet email
 - Mail between users on multi-user machines had existed since the 1960s
 - With the advent of ARPANET (1970s), mail began to be transferred over networks
 - Initially, mail transfer was performed over FTP
 - Sites connected intermittently (usually via dialup serial link) transferred mail using UUCP
 - UUCP allows for transfer of files over serial links
 - ARPANET-UUCP gateways allowed much faster transfers

History con't

- Without direct connections between all sites, and no way of determining how best to get a message there, senders specified the routing of their messages
 - The infamous bang-paths: ...foovax!barbox! quuxnet!me, where presumably everyone could find a path to foovax
 - People would often specify paths from multiple hosts: ...{ucbvax,menlo70}!barbox!me
 - Classic example: ...{decvax,philabs}!mcvax! moskvax!kremvax!chernenko

- The modern era: SMTP
 - The protocol is defined in RFC 2821, the message format in RFC 2822
 - Messages are transferred in a simple text format: headers, newline, body
 - Headers: From, To, Subject, Date, ...
 - Content of the body may not be plain text in modern usage – see the MIME RFCs
 - Note that the information in the headers is NOT used in routing the message

SMTP con't

- Routing an email message
 - The destination of an email is determined by its envelope recipient, specified to the server on sending
 - The sending mail server looks for an MX record for the destination site to determine who to talk to
 - MX records are tried in order of the priority given
 - If no MX record is present, an MX record pointing to the name itself with priority 0 is assumed
 - If the message doesn't get there, the envelope sender is returned a bounce

- SMTP con't
 - A typical SMTP conversation
 - HELO hostname or EHLO hostname
 - MAIL FROM: <envelope-sender@site>
 - RCPT T0: <envelope-recipient@othersite>
 - DATA
 - QUIT
 - The server replies with the traditional 1xx-5xx codes

- Problems with SMTP
 - No identity verification anyone can send email claiming to be anyone else
 - In fact, most mailing list software depends on this behavior
 - One of the enabling factors behind spam, blowback bounces, etc.
 - Legacy compatibility causes problems
 - Long timeouts, retries, etc. for unreliable networks lead to messages spending in mail queues before bouncing
 - Bang paths still (theoretically) supported!

Mail Transfer Agents

- sendmail
 - The original Internet mailer, an extension by Eric Allman of the original UUCP-era mailers
 - Complex configuration, but can do (literally) anything
 - Inefficient, slow, terrible security record
 - Includes a nice mail filtering architecture in newer versions (milter)
 - Not recommended for new installs unless you have special needs

MTAs con't

- qmail
 - Dan Bernstein's attempt at a secure mailer
 - Multi-binary, privilege-separated architecture provides better security and faster mail delivery
 - Introduced maildir mail storage (one file per email); often faster than mbox (one file per mail spool)
 - Now public-domain and "maintained" by a (vocal) community; has some quirky behaviors
 - Not recommended for new installs

MTAs con't

- Exim
 - Developed at Cambridge as a fast Internet mailer
 - Single-binary architecture sometimes criticized, but has reasonable security record
 - Extremely fast in environments where most mail is deliverable immediately, but has poor queue management
 - Can embed Perl to process email

- MTAs con't
 - Postfix
 - Wietse Venema's attempt at a secure mailer
 - Multi-binary, privilege-separated architecture
 - Very fast, simple configuration
 - Some processing isn't possible inside Postfix, but external hooks are possible
 - Such filtering setups can be a bit complex
 - Sendmail milters are supported to some degree in newer versions